

# Multitask Training and Instruction-Based Fine-Tuning Flan T5/ T0



Abdelaziz Guelfane Imane Meziany Malek Bouhadida Mohammed El Barhichi Yousra Yakhou



- Pre-2018 : Task-Specific Training
  - RNNs, LSTMs, GRUs were trained on specific tasks.
  - Even Seq2Seq models with Attention (Transformer architecture) involved training a separate model for each task.



- Pre-2018 : Task-Specific Training
  - RNNs, LSTMs, GRUs were trained on specific tasks.
  - Even Seq2Seq models with Attention (Transformer architecture) involved training a separate model for each task.

#### • 2019 : Text-To-Text framework

- T5 reformulates any NLP task to a text generation problem.
- Example : Instead of predicting a class label, the model generates the label as text (e.g., "positive sentiment").
- However, fine-tuning on specific tasks was still needed for optimal performance.



- Pre-2018 : Task-Specific Training
  - RNNs, LSTMs, GRUs were trained on specific tasks
  - Even Seq2Seq models with Attention (Transformer architecture) involved training a separate model for each task.

#### • 2019 : Text-To-Text framework

- T5 reformulates any NLP task to a text generation problem.
- Example : Instead of predicting a class label, the model generates the label as text (e.g., "positive sentiment").
- However, fine-tuning on specific tasks was still needed for optimal performance.

#### • 2020-2022 : Instruction Tuning on Multi-tasks

- T0: Sanh, V., A. Webson, C. Raffel, et al. "Multitask Prompted Training Enables Zero-Shot Task Generalization." In ICLR, 2022.
- Flan-T5: Chung, H.W., L. Hou, S. Longpre, et al. "Scaling Instruction-Finetuned Language Models." In CoRR, 2022.



- Pre-2018 : Task-Specific Training
  - RNNs, LSTMs, GRUs were trained on specific tasks
  - Even Seq2Seq models with Attention (Transformer architecture) involved training a separate model for each task.

#### • 2019 : Text-To-Text framework

- T5 reformulates any NLP task to a text generation problem.
- Example : Instead of predicting a class label, the model generates the label as text (e.g., "positive sentiment").
- However, fine-tuning on specific tasks was still needed for optimal performance.

#### • 2020-2022: Instruction Tuning on Multi-tasks

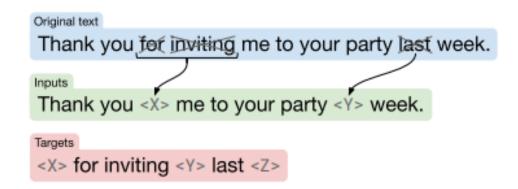
#### **Better Generalization?**

- T0: Sanh, V., A. Webson, C. Raffel, et al. "Multitask Prompted Training Enables Zero-Shot Task Generalization." In ICLR, 2022.
- Flan-T5: Chung, H.W., L. Hou, S. Longpre, et al. "Scaling Instruction-Finetuned Language Models." In CoRR, 2022.

#### **T5: Text-to-Text Transfer Transformer**



- T5 reformulates any NLP task as a sequence-to-sequence generation problem.
- It is based on an **Encoder-Decoder** architecture and trained using span corruption.
- It was pretained on C4 (Colossal Clean Crawled Corpus).



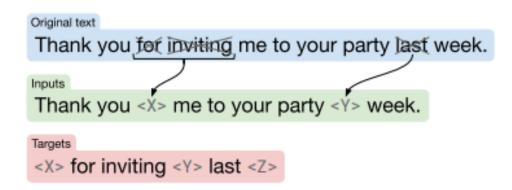
- **T5+LM** is the T5 version designed for language modeling tasks.
- It is T5 trained on 100B additional tokens from C4 on a causal language modeling objective.

The model predicts the next token in a sequence, given all prior tokens.

#### **T5: Text-to-Text Transfer Transformer**



- **T5** reformulates any NLP task as a sequence-to-sequence generation problem.
- It is based on an **Encoder-Decoder** architecture and trained using span corruption.
- It was pretained on C4 (Colossal Clean Crawled Corpus).



- T5+LM is the T5 version designed for language modeling tasks.
- It is T5 trained on 100B additional tokens from C4 on a causal language modeling objective.

The model predicts the next token in a sequence, given all prior tokens.

• T5-LM, and LLMs in general, attain reasonable zero-shot generalization on a diverse set of tasks.

#### Hypothesis:



Consequence of implicit multitask learning in language models' pretraining.



#### **Implicit multitask learning**

Many websites contain lists of trivia Q&As

>>>> Supervised training data for the task of closedbook question answering.

1. What does "www" stand for in a website browser?

Answer: World Wide Web

2. How long is an Olympic swimming pool (in meters)?

Answer: 50 meters

3. What countries made up the original Axis powers in World War II?

Answer: Germany, Italy, and Japan



#### **Implicit multitask learning**

- Many websites contain lists of trivia Q&As
- >>>> Supervised training data for the task of closedbook question answering.

1. What does "www" stand for in a website browser?

Answer: World Wide Web

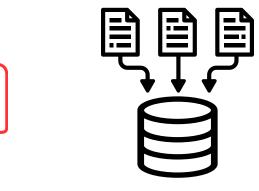
2. How long is an Olympic swimming pool (in meters)?

Answer: 50 meters

3. What countries made up the original Axis powers in World War II?

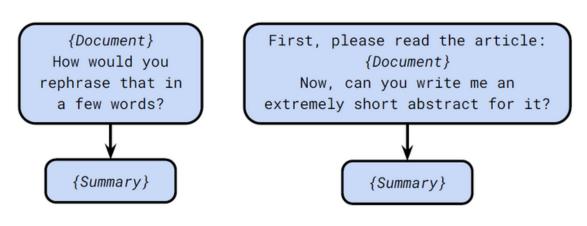
Answer: Germany, Italy, and Japan

### Limitations..



Requires a sufficiently large model + larger corpus of data

2



Sensitive to the wording of prompts



#### **Implicit multitask learning**

- Many websites contain lists of trivia Q&As
- >>>> Supervised training data for the task of closedbook question answering.



1. What does "www" stand for in a website browser?

Answer: World Wide Web

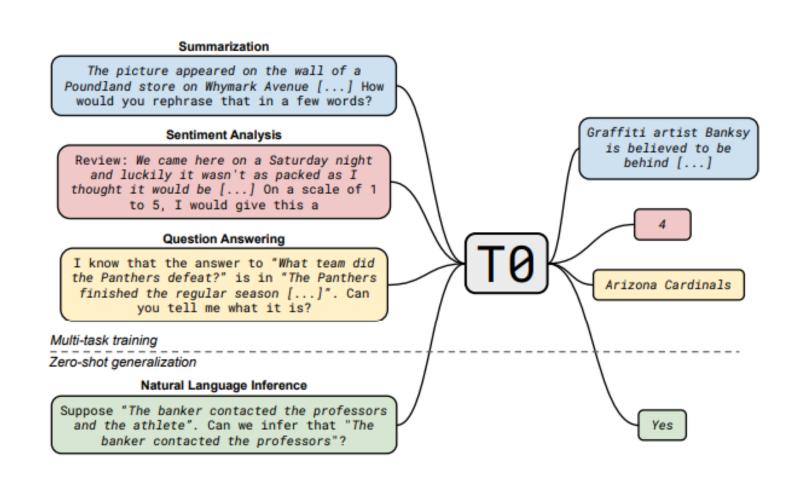
2. How long is an Olympic swimming pool (in meters)?

Answer: 50 meters

3. What countries made up the original Axis powers in World War II?

Answer: Germany, Italy, and Japan

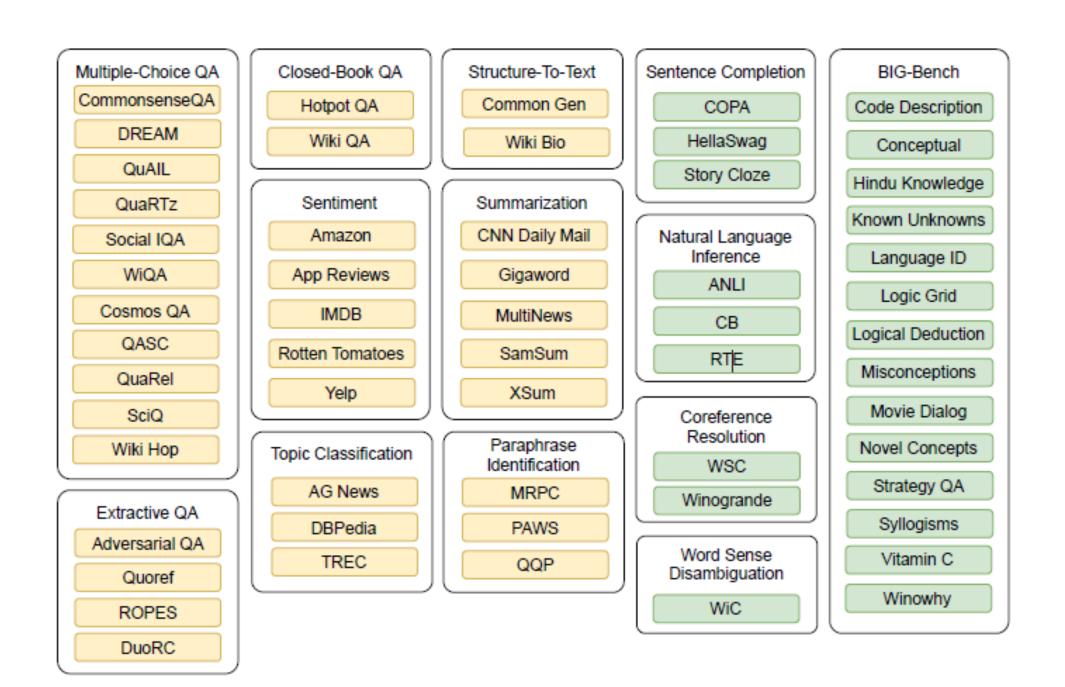
#### **Explicit multitask learning**





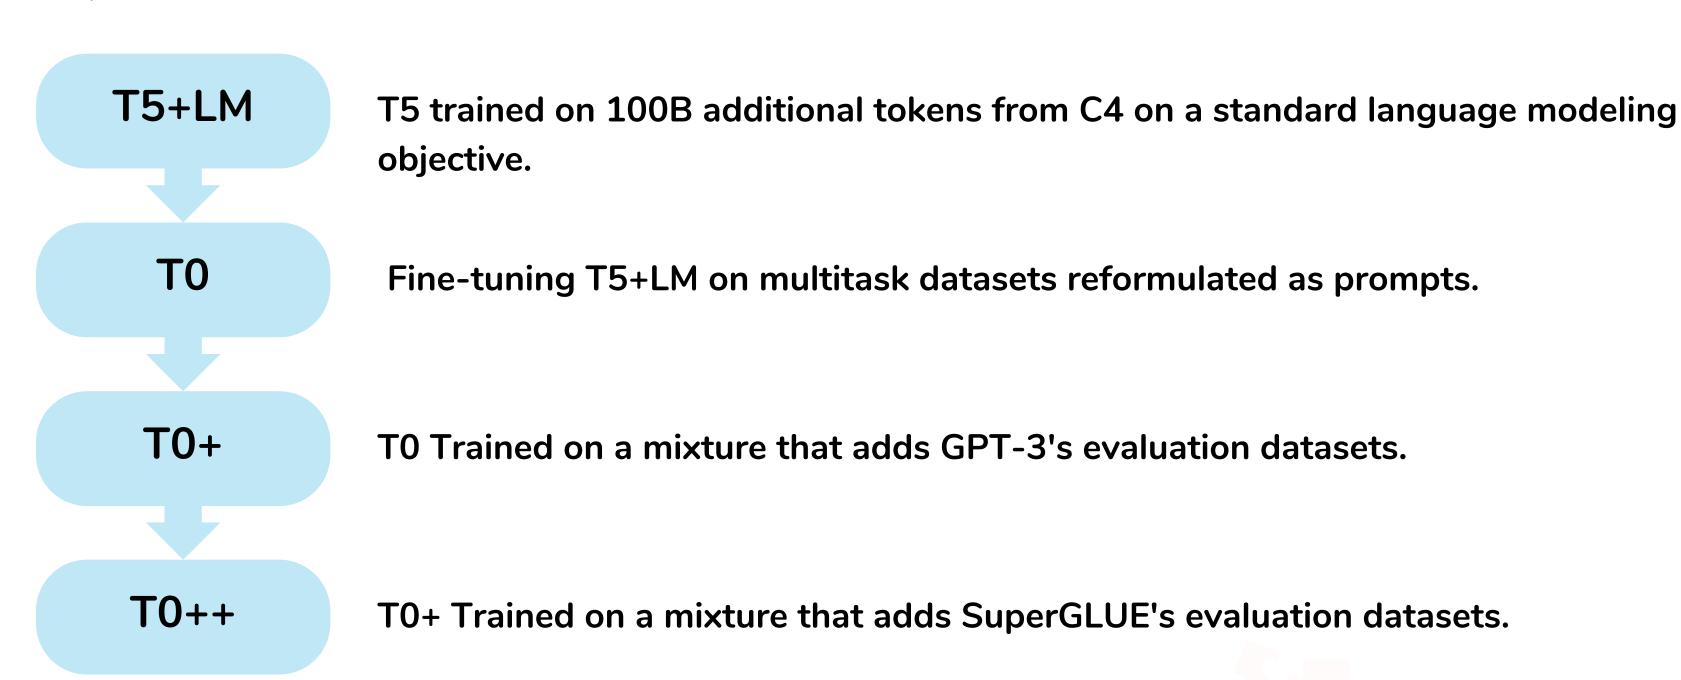
• T0 datasets and task taxonomy (T0+ and T0++ are trained on additional datasets).

- 12 tasks and 62 datasets with publicly contributed prompts.
- Yellow datasets are in the training mixture.
- Green datasets are held out and represent tasks that were not seen during training.





• A quick overview of the models :

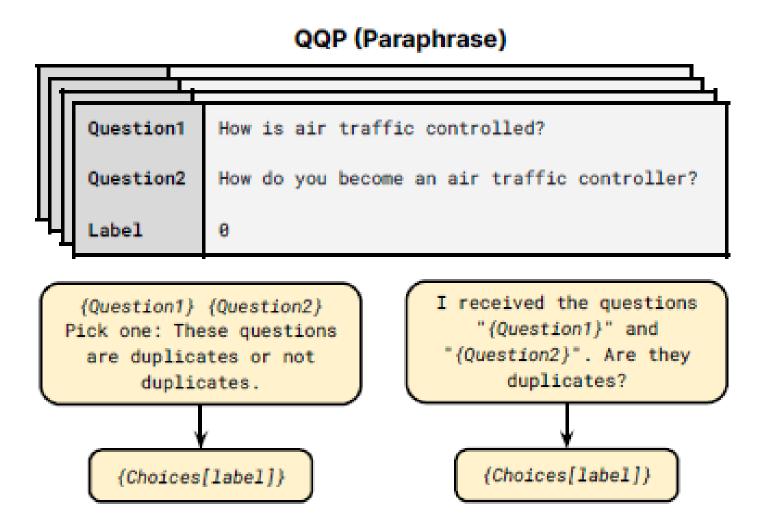








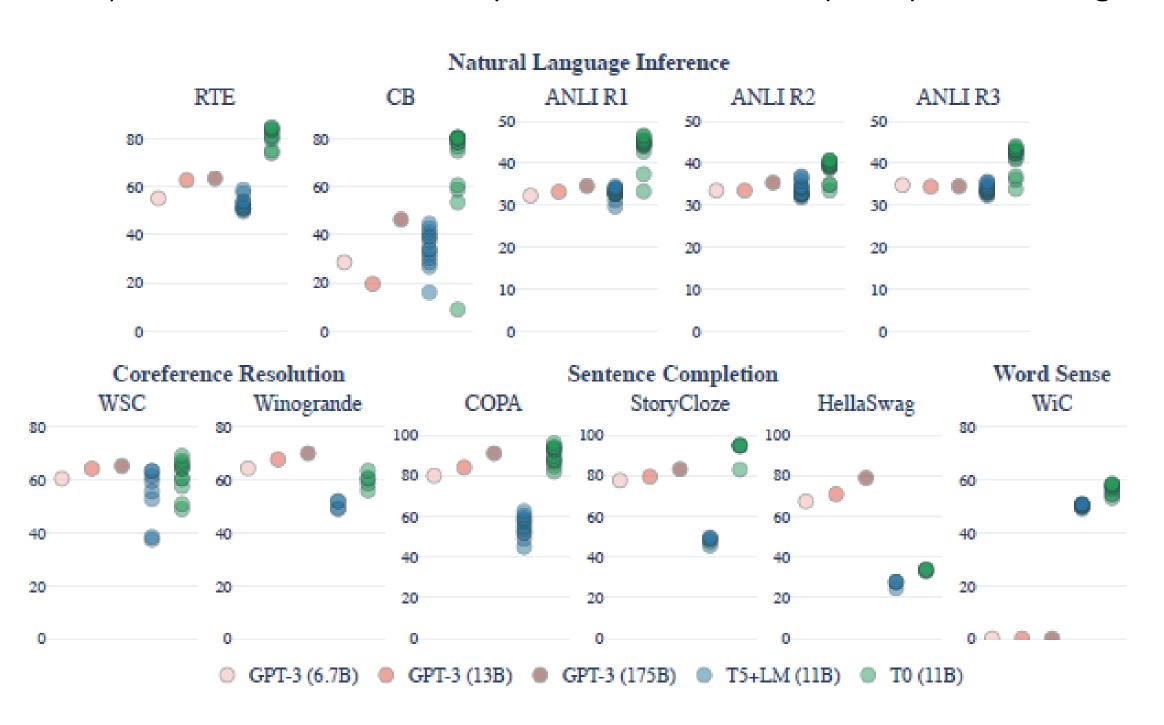
Models learn to understand the prompts as task instructions which help them generalize to heldout tasks







- Results for T0 task generalization experiments compared to GPT-3.
- T5+LM (baseline model) is the same as T0 except without multitask prompted training.




Better Generalization!



• Results for a subset of BIG-bench models compared to T0, T0+ and T0++.

# Better Generalization!





• So far, we have seen that..




>>>>> Multitask prompted training improve generalization to held-out tasks

• Now,

>>>>> Does training on a wider range of prompts improve robustness to prompt wording?



• Zero-shot performance of T0 and T5+LM when increasing number of training prompts per dataset.





What happens if we scale even further?

- Increasing the **number of tasks** and **data diversity**?
- Leveraging larger models?
- Incorporating **reasoning capabilities** like Chain-of-Thought (CoT)?



#### What happens if we scale even further?

- Increasing the number of tasks and data diversity?
- Leveraging larger models?
- Incorporating **reasoning capabilities** like Chain-of-Thought (CoT)?

#### Scaling Instruction-Finetuned Language Models

Hyung Won Chung\* Le Hou\* Shayne Longpre\* Barret Zoph\* Yi Tay\*
William Fedus\* Yunxuan Li Xuezhi Wang Mostafa Dehghani Siddhartha Brahma
Albert Webson Shixiang Shane Gu Zhuyun Dai Mirac Suzgun Xinyun Chen
Aakanksha Chowdhery Alex Castro-Ros Marie Pellat Kevin Robinson
Dasha Valter Sharan Narang Gaurav Mishra Adams Yu Vincent Zhao
Yanping Huang Andrew Dai Hongkun Yu Slav Petrov Ed H. Chi
Jeff Dean Jacob Devlin Adam Roberts Denny Zhou Quoc V. Le

Google

#### Abstract

Finetuning language models on a collection of datasets phrased as instructions has been shown to improve model performance and generalization to unseen tasks. In this paper we explore instruction finetuning with a particular focus on (1) scaling the number of tasks, (2) scaling the model size, and (3) finetuning on chain-of-thought data. We find that instruction finetuning with the above aspects dramatically improves performance on a variety of model classes (PaLM, TS, U-PaLM), prompting setups (zero-shot, few-shot, CoT), and evaluation benchmarks (MMLU, BBH, TyDiQA, MGSM, open-ended generation, RealToxicityPrompts). For instance, Flan-PaLM 540B instruction-finetuned on 1.8K tasks outperforms PaLM 540B by a large margin (+9.4% on average). Flan-PaLM 540B achieves state-of-the-art performance on several benchmarks, such as 75.2% on five-shot MMLU. We also publicly release Flan-T5 checkpoints, which achieve strong few-shot performance even compared to much larger models, such as PaLM 62B. Overall, instruction finetuning is a general method for improving the performance and usability of pretrained language models.




Figure 1: We finetune various language models on 1.8K tasks phrased as instructions, and evaluate them on unseen tasks. We finetune both with and without exemplars (i.e., zero-shot and few-shot) and with and without chain-of-thought, enabling generalization across a range of evaluation scenarios.

9


1

<sup>\*</sup>Equal contribution. Correspondence: lehou@google.com.


<sup>&</sup>lt;sup>†</sup>Core contributor

<sup>&</sup>lt;sup>1</sup>Public checkpoints: https://github.com/google-research/t5x/blob/main/docs/models.md#flan-t5-checkpoints.








# Scaling Tasks and Datasets



























• Dataset used:

• Scaling Finetuning Tasks:

473 datasets
146 task categories
1 836 Tasks!

#### Finetuning tasks

#### TO-SF

Commonsense reasoning
Question generation
Closed-book QA
Adversarial QA
Extractive QA
Title/context generation
Topic classification
Struct-to-text

55 Datasets, 14 Categories, 193 Tasks

#### Muffin

Natural language inference Closed-book QA
Code instruction gen. Conversational QA
Program synthesis Code repair
Dialog context generation ...

69 Datasets, 27 Categories, 80 Tasks

#### CoT (Reasoning)

Arithmetic reasoning Explanation generation
Commonsense Reasoning Sentence composition
Implicit reasoning ...

9 Datasets, 1 Category, 9 Tasks

#### Natural Instructions v2

Cause effect classification
Commonsense reasoning
Named entity recognition
Toxic language detection
Question answering
Question generation
Program execution
Text categorization

---

372 Datasets, 108 Categories, 1554 Tasks

- A <u>Dataset</u> is an original data source (e.g. SQuAD).
- A <u>Task Category</u> is unique task setup (e.g. the SQuAD dataset is configurable for multiple task categories such as extractive question answering, query generation, and context generation).
- A <u>Task</u> is a unique <dataset, task category> pair, with any number of templates which preserve the task category (e.g. query generation on the SQuAD dataset.)



Dataset used :

Scaling Finetuning Tasks:

473 datasets 146 task categories 1 836 Tasks!

• Evaluation :

#### Finetuning tasks

Muffin

Closed-book QA

Code repair

Conversational QA

#### TO-SF

Commonsense reasoning Question generation Closed-book QA Adversarial QA Extractive QA Title/context generation Topic classification

55 Datasets, 14 Categories,

193 Tasks

Struct-to-text

Natural language inference Code instruction gen. Program synthesis

Dialog context generation

69 Datasets, 27 Categories, 80 Tasks

CoT (Reasoning)

Arithmetic reasoning Explanation generation Commonsense Reasoning Sentence composition Implicit reasoning

9 Datasets, 1 Category, 9 Tasks

#### Natural Instructions v2

Cause effect classification Commonsense reasoning Named entity recognition Toxic language detection Question answering Question generation Program execution Text categorization

372 Datasets, 108 Categories, 1554 Tasks

- A Dataset is an original data source (e.g. SQuAD).
- A Task Category is unique task setup (e.g. the SQuAD dataset is configurable for multiple task categories such as extractive question answering, query generation, and context generation).
- A Task is a unique <dataset, task category> pair, with any number of templates which preserve the task category (e.g. query generation on the SQuAD dataset.)

#### Held-out tasks

#### MMLU

Abstract algebra College medicine Professional law

Sociology Philosophy

57 tasks

#### **BBH**

Boolean expressions Navigate Tracking shuffled objects Word sorting Dyck languages

27 tasks

#### **TvDiQA**

Information seeking QA

8 languages

#### MGSM

Grade school math problems

10 languages



• Tasks & Instructions variety:

# Finetuning includes more diverse tasks :

- zero-shot
- few-shot
- chain-of-thought



• Tasks & Instructions variety:

Finetuning includes more diverse tasks:

- zero-shot
- few-shot
- chain-of-thought

Without chain-of-thought Answer the following yes/no question. Instruction Can you write a whole exemplars Haiku in a single tweet?

without

Without chain-of-thought



• Tasks & Instructions variety:

Finetuning includes more diverse tasks:

- zero-shot
- few-shot
- chain-of-thought

Answer the following Instruction yes/no question. without Can you write a whole exemplars Haiku in a single tweet? Q: Answer the following yes/no question. Could a dandelion suffer from hepatitis? Instruction A: no with exemplars Q: Answer the following yes/no question. Can you write a whole Haiku

in a single tweet?



Tasks & Instructions variety :

Finetuning includes more diverse tasks:

- zero-shot
- few-shot
- chain-of-thought

Instruction without exemplars

Instruction

with exemplars

Answer the following yes/no question.

Can you write a whole Haiku in a single tweet?

Without chain-of-thought

Q: Answer the following yes/no question.
Could a dandelion suffer from hepatitis?
A: no
Q: Answer the following yes/no question.
Can you write a whole Haiku in a single tweet?
A:

With chain-of-thought

Answer the following yes/no question by reasoning step-by-step.

Can you write a whole Haiku in a single tweet? A haiku is a japanese three-line poem. That is short enough to fit in 280 characters. The answer is yes.

Q: Answer the following yes/no question by reasoning step-by-step.

Could a dandelion suffer from hepatitis?

A: Hepatitis only affects organisms with livers.

Dandelions don't have a liver. The answer is no.

Q: Answer the following yes/no question by reasoning step-by-step. Can you write a whole Haiku in a single tweet? A haiku is a japanese three-line poem. That is short enough to fit in 280 characters. The answer is yes.



#### • Models' Sizes:

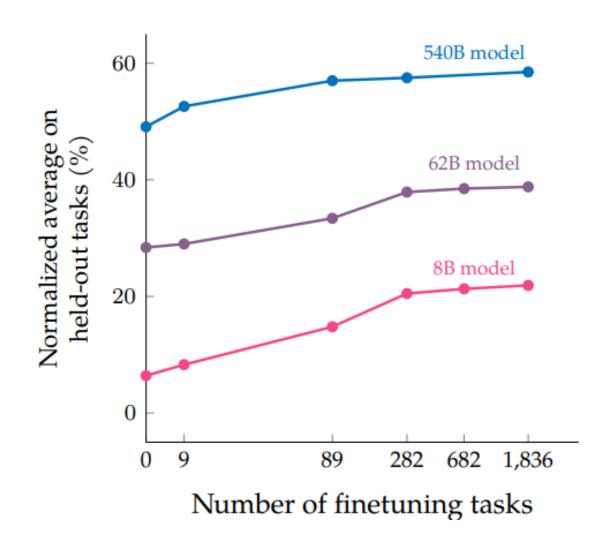
Params	Model	Architecture	Pre-training Objective			
80M	Flan-T5-Small	encoder-decoder	span corruption			
250M	Flan-T5-Base	encoder-decoder	span corruption			
780M	Flan-T5-Large	encoder-decoder	span corruption			
3B	Flan-T5-XL	encoder-decoder	span corruption			
11B	Flan-T5-XXL	encoder-decoder	span corruption			
8B	Flan-PaLM	decoder-only	causal LM			
62B	Flan-PaLM	decoder-only	causal LM			
540B	Flan-PaLM	decoder-only	causal LM			
62B	Flan-cont-PaLM	decoder-only	causal LM			
540B	Flan-U-PaLM	decoder-only	prefix LM + span corruption			



#### • Results:

#### Performance improves on:

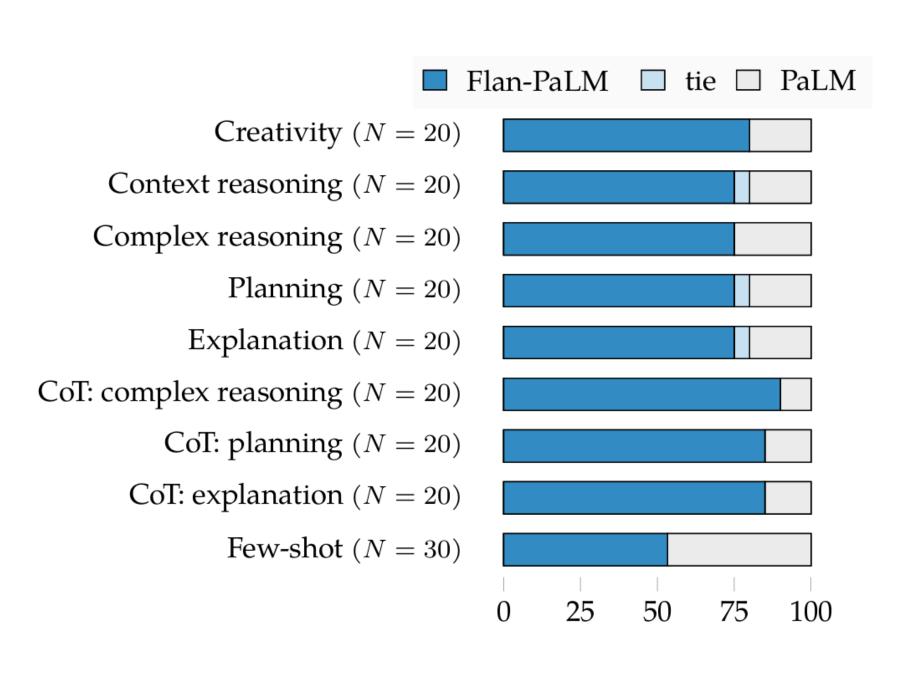
- different model classes
- prompting setups
- evaluation benchmarks

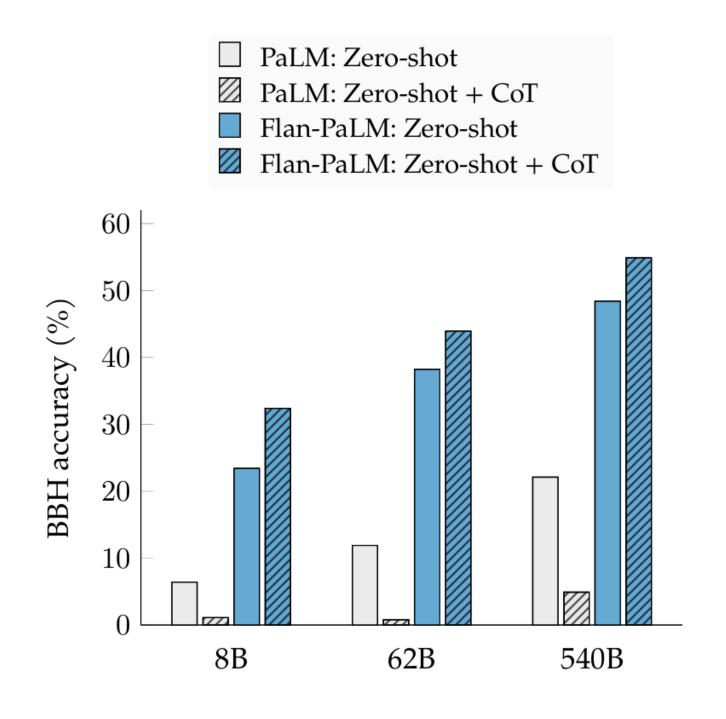


			MMLU		BBH		TyDiQA	MGSM
Params	Model	Norm. avg.	Direct	CoT	Direct	CoT	Direct	CoT
80M	T5-Small	-9.2	26.7	5.6	27.0	7.2	0.0	0.4
	Flan-T5-Small	-3.1 (+6.1)	28.7	12.1	29.1	19.2	1.1	0.2
250M	T5-Base	-5.1	25.7	14.5	27.8	14.6	0.0	0.5
	Flan-T5-Base	6.5 (+11.6)	35.9	33.7	31.3	27.9	4.1	0.4
780M	T5-Large	-5.0	25.1	15.0	27.7	16.1	0.0	0.3
	Flan-T5-Large	13.8 (+18.8)	45.1	40.5	37.5	31.5	12.3	0.7
3B	T5-XL	-4.1	25.7	14.5	27.4	19.2	0.0	0.8
	Flan-T5-XL	19.1 (+23.2)	52.4	45.5	41.0	35.2	16.6	1.9
11B	T5-XXL	-2.9	25.9	18.7	29.5	19.3	0.0	1.0
	Flan-T5-XXL	23.7 (+26.6)	55.1	48.6	45.3	41.4	19.0	4.9
8B	PaLM	6.4	24.3	24.1	30.8	30.1	25.0	3.4
	Flan-PaLM	21.9 (+15.5)	49.3	41.3	36.4	31.1	47.5	8.2
62B	PaLM	28.4	55.1	49.0	37.4	43.0	40.5	18.2
	Flan-PaLM	38.8 (+10.4)	59.6	56.9	47.5	44.9	58.7	28.5
540B	PaLM	49.1	71.3	62.9	49.1	63.7	52.9	45.9
	Flan-PaLM	58.4 (+9.3)	73.5	70.9	57.9	66.3	67.8	57.0
62B	cont-PaLM	38.1	61.2	57.6	41.7	53.1	45.7	32.0
	Flan-cont-PaLM	46.7 (+8.6)	66.1	62.0	51.0	53.3	62.7	40.3
540B	U-PaLM	50.2	71.5	64.0	49.2	62.4	54.6	49.9
,	Flan-U-PaLM	59.1 (+8.9)	74.1	69.8	59.3	64.9	68.3	60.4



#### • Results:





Instruction Tuning Improves Human Usability

Finetuning with Chain-of-Thought Data Unlocks Zero-Shot Reasoning

#### References



• Chung, H. W., Hou, L., Longpre, S., Zoph, B., Tay, Y., Fedus, W., Li, Y., Wang, X., Dehghani, M., Brahma, S., Webson, A., Gu, S. S., Dai, Z., Suzgun, M., Chen, X., Chowdhery, A., Castro-Ros, A., Pellat, M., Robinson, K., . . Wei, J. (2022, 20 octobre). Scaling Instruction-Finetuned Language Models. arXiv.org. https://arxiv.org/abs/2210.11416

Sanh, V., Webson, A., Raffel, C., Bach, S. H., Sutawika, L., Alyafeai, Z., Chaffin, A., Stiegler, A., Scao, T. L., Raja, A., Dey, M., Bari, M. S., Xu, C., Thakker, U., Sharma, S. S., Szczechla, E., Kim, T., Chhablani, G., Nayak, N.,... Rush, A. M. (2021, 15 octobre).
 Multitask prompted training enables Zero-Shot task generalization. arXiv.org. https://arxiv.org/abs/2110.08207



# THANK YOU



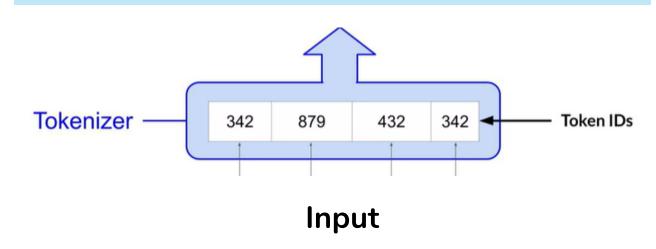
## **Appendix**

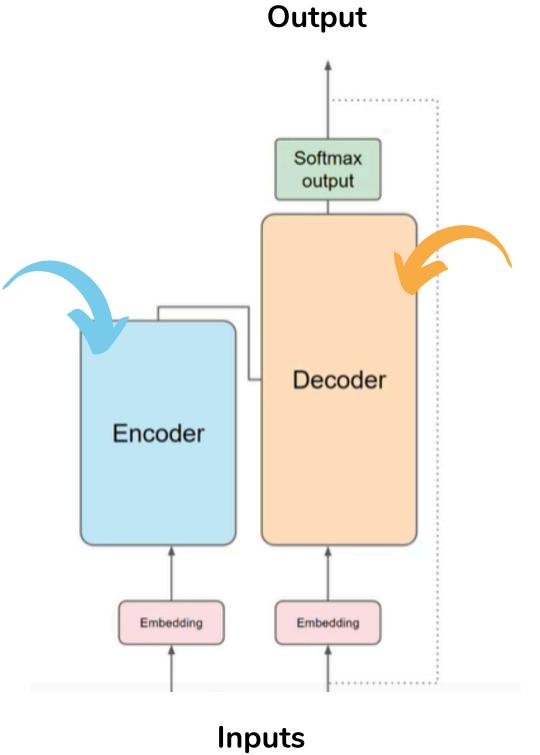


#### sequence vectors:

encodes information about both the element itself and its context within the sequence.

process the input sequence and create contextualized representations for each element (word or token) in the sequence





the final representation vectors

generate an output
sequence based on the
contextualized
representations
provided by the
encoder