


Scaling Laws for Neural Language Models (2020)

Jared Kaplan, Sam McCandlish et Al.
from OpenAI




La vérité sur les
Scaling Laws.

Context

- **Language as a Natural Domain for AI:**
 - Most **reasoning tasks** can be resolved with language.
 - The abundance of **text data worldwide** permits powerful learning.
- **Rapid Progress in Deep Learning for LLM:**
 - Models closer to **human-level performance** on many tasks (Transformers, BERT ...)
 - Models now excel in tasks like composing **coherent**, **multi-paragraph prompted text samples**.

The Rise and Rise of A.I. Large Language Models (LLMs) & their associated bots like ChatGPT



Kaplan et
al.

Hoffmann et al.

Art1 - Scaling Laws for Neural Language Models

Why this article?

Goal:

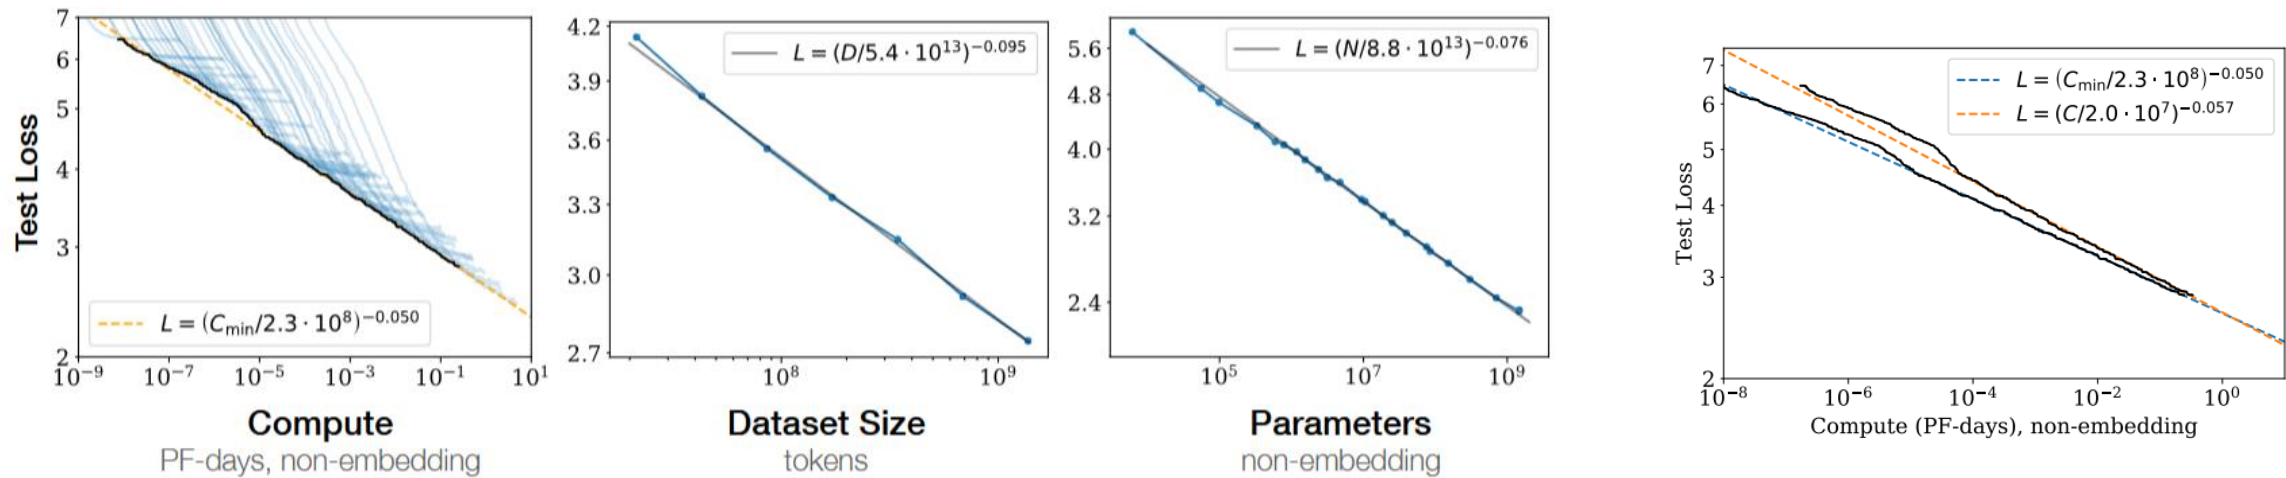
- Empirically investigate how **language modeling performance** (ie the Loss function) is influenced by **various factors**
- Based on this analysis, determine the optimal **trade-off** between **model size** and **dataset size** for a given **training compute budget**.

Results:

- Influence of **factors**:
 - Training compute (C)
 - Dataset size (D)
 - Model size (N)
 - Model shapes (Width & Depth)
 - Batch size (B)
 - Number of steps (S)
- Loss follows power laws
- On very large models, it is more efficient to stop training early
- $D \propto N^{0.74}$ (False: Chinchilla) should evolve in tandem
- Existence of optimal batch size
- Transfer incurs only a constant penalty
- Sample efficiency

Art1 - Scaling Laws for Neural Language Models

Main results – Empirical laws



With X as the factor, X_c as the constant (meaningless) and α as the scale factor :


$$L(X) = \left(\frac{X_c}{X}\right)^{\alpha_X}$$

Power Law	Scale (tokenization-dependent)
$\alpha_N = 0.076$	$N_c = 8.8 \times 10^{13}$ params (non-embed)
$\alpha_D = 0.095$	$D_c = 5.4 \times 10^{13}$ tokens
$\alpha_C = 0.057$	$C_c = 1.6 \times 10^7$ PF-days
$\alpha_C^{\min} = 0.050$	$C_c^{\min} = 3.1 \times 10^8$ PF-days
$\alpha_B = 0.21$	$B_* = 2.1 \times 10^8$ tokens
$\alpha_S = 0.76$	$S_c = 2.1 \times 10^3$ steps

Art1 - Scaling Laws for Neural Language Models

Main results - Model shape dependency & Sample Efficiency

Performances are **weakly** affected by **model shape** for a Transformer.




Feed-Forward Width

Depth

Number of Attention heads

d_{model} is the dimensionality of the model's hidden states

S_{min} : Estimated **steps** to obtain a **given compute**



The more the **size grows**, the more the model is **sample efficient**

For a **given compute**, you can go with a **big model** with **early stopping**

Art1 - Scaling Laws for Neural Language Models

Experiment

- **Training on WebText2 (dataset)**

- $n_{\text{vocab}} = 50257$
- Loss : cross-entropy over 1024 token context

- **Training factors:**

- Variation of non-embedding [parameters number](#)
768M \rightarrow 1500M parameters
- Adam optimizer with a [fixed number of training tokens](#) (most of runs)
2.5x10⁵ steps with batches of 512 sequences of 1024 tokens
- Learning rate schedule : linear warmup followed by a cosine decay





Figure 2 – some of the training runs

Art1 - Scaling Laws for Neural Language Models

Experiment


$$L(N, D) = \left[\left(\frac{N_c}{N} \right)^{\frac{\alpha_N}{\alpha_D}} + \frac{D_c}{D} \right]^{\alpha_D}$$



- **Vocabulary size changes :**
 - Changes in vocabulary size or tokenization should scale the loss proportionally, so $L(N, D)$ must support this scaling naturally
- **Asymptotic behavior of loss :**
 - As $N \rightarrow \infty$ (with fixed D), the loss should approach $L(D)$; as $D \rightarrow \infty$ (with fixed N), it should approach $L(N)$
- **Series expansion at large D :**
 - $L(N, D)$ should allow a series expansion in $1/D$, though this principle has weaker theoretical backing.

Art1 - Scaling Laws for Neural Language Models

Contradiction




To keep **underfitting under control** :

$$D \propto N^{0.74} \propto C_{\min}^{0.54}$$

While at compute-efficient training : (for $C = 2 \cdot C_{\min}$) :

$$D(C_{\min}) = \frac{2C_{\min}}{6N(C_{\min})} \propto C_{\min}^{0.26}$$



Optimal values for the **minimal loss** theoretically reachable

$$C^* \sim 10^4 \text{ PF-Days} \quad N^* \sim 10^{12} \text{ parameters,}$$

$$D^* \sim 10^{12} \text{ tokens, } L^* \sim 1.7 \text{ nats/token}$$



Training Compute-Optimal Large Language Models (2022)

Jordan Hoffmann, Sebastian Borgeaud, Arthur Mensch

From DeepMind

J. Kaplan vous a
menti. Ma réponse

Art2 - Training Compute-Optimal Large Language Models (2022)

Goal

- Goal: same as in the previous article:

best model size / dataset size trade-off

- Differences:

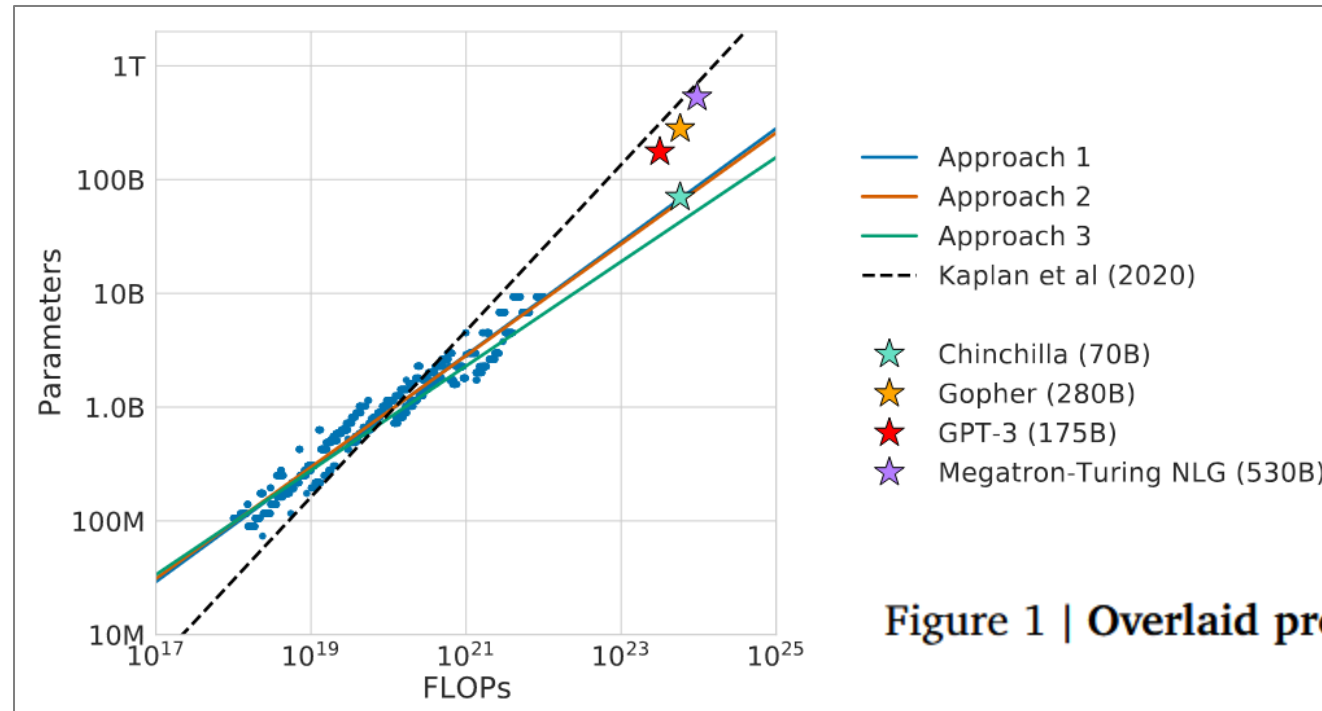
- More model parameters than before:
 - art1: <100M parameters
 - art2: >500M parameters
- Different dataset size and learning rate
- Approach as an optimization problem: **Minimizing $L(N, D)$** under the constraint **$FLOPs(N, D) = C$**

$$N_{opt}(C), D_{opt}(C) = \underset{N, D \text{ s.t. } FLOPs(N, D) = C}{\operatorname{argmin}} L(N, D).$$

Art2 - Training Compute-Optimal Large Language Models (2022)

Main results – new scaling laws

- A new scaling law:




- Most of the models on market have **too much parameters compared to the size of there dataset**
- Disagreement with art1 : **Model size x2 → Dataset x2**

Art2 - Training Compute-Optimal Large Language Models (2022)

3 approaches


Approach 1

- Varying training sequences (4 different sizes for each model)
- Fixed model sizes (from 70M to 10B)



Approach 2 - IsoFLOPS

- Varying model size (up to 16B)
- Fixed FLOPS count (9 different)




Approach 3 – Fitting law

- Fitting the law:

$$\hat{L}(N, D) \triangleq E + \frac{A}{N^\alpha} + \frac{B}{D^\beta}$$

- Using the loss :

$$\min_{A, B, E, \alpha, \beta} \sum_{\text{Runs } i} \text{Huber}_\delta \left(\log \hat{L}(N_i, D_i) - \log L_i \right)$$



Art2 - Training Compute-Optimal Large Language Models (2022)

Result : Chinchilla

- A small model compare to others (70B)

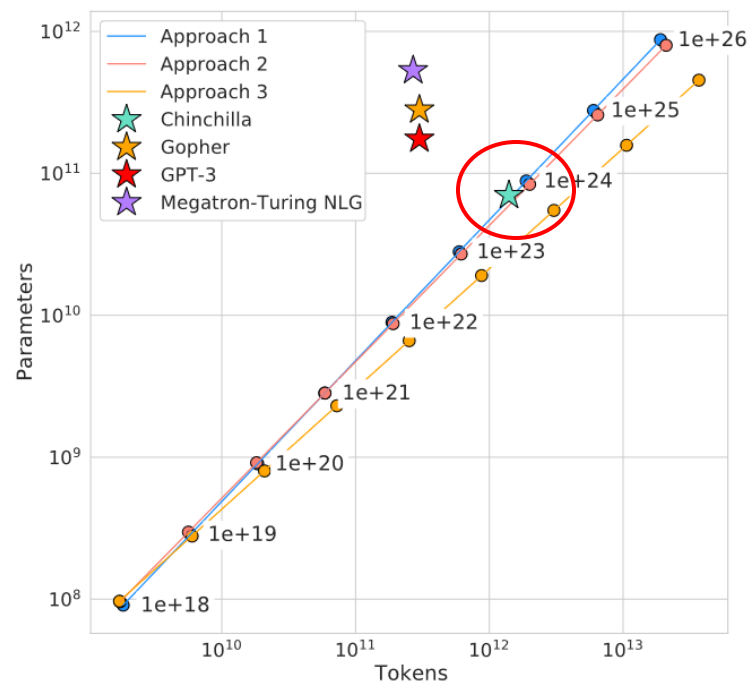


Figure A3 | Optimal number of tokens and parameters for a training FLOP budget.

- Outperform Gopher on most tasks




Figure 6 | MMLU results compared to Gopher We find that Chinchilla outperforms Gopher by 7.6% on average (see Table 6) in addition to performing better on 51/57 individual tasks, the same on


Random	25.0%
Average human rater	34.5%
GPT-3 5-shot	43.9%
Gopher 5-shot	60.0%
Chinchilla 5-shot	67.6%
Average human expert performance	89.8%
June 2022 Forecast	57.1%
June 2023 Forecast	63.4%

Table 6 | Massive Multitask Language Understanding (MMLU).

Conclusion

- After Kaplan:
 - Tendency to increase more the model size than the dataset size
- After Chinchilla:
 - Slowed the model increase pace
 - Allocate more computation on dataset size
 - Tendency to increase the model size as much as the dataset size

The Rise and Rise of A.I. Large Language Models (LLMs) & their associated bots like ChatGPT



Kaplan et al.

Hoffmann et al.